

# Lane Segmentation Week 1

# 主要内容

- 计算机视觉学习提速绝招
- 项目概述
- ResNet
- ResNeSt

# 朱利明



人工智能资深讲师，华为云AI社区达人

2002年毕业于中科院研究生院，近二十年软件和算法研发经验，主要研究领域为自动驾驶、人体姿态识别、3D计算机视觉，具有京东、华为等大厂AI资源，在人工智能课程研发和教学有丰富经验，已指导近千位同学进入计算机视觉领域。

微信：aiking2018

# 入行的最短时间

- 六个月-三个月

# 学习提速绝招

- 明确目标
- 确定最短学习路径
- 选择正确的参考书
- 必学必会

# 明确学习目标

- 求职
- 毕设
- 提高水平

# 确定最短学习路径

- AI是一个开放领域，没有明确的考试大纲
- 有所学有所不学

# 预备知识

- Python
- Machine Learning
- CNN and Image Classification
- Tensorflow or Pytorch
- English
- Math



# 你听到的

- 数学要求高（高数、线性代数、概率论）
- 要懂机器学习
- 要精通Python
- 要会Tensorflow/Pytorch
- 要会OpenCV.....
- 要会C++


# 你看到的

$$\begin{aligned} & \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[ (x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 \right] \\ & + \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[ \left( \sqrt{w_i} - \sqrt{\hat{w}_i} \right)^2 + \left( \sqrt{h_i} - \sqrt{\hat{h}_i} \right)^2 \right] \\ & + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left( C_i - \hat{C}_i \right)^2 \\ & + \lambda_{\text{noobj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{noobj}} \left( C_i - \hat{C}_i \right)^2 \\ & + \sum_{i=0}^{S^2} \mathbb{1}_i^{\text{obj}} \sum_{c \in \text{classes}} (p_i(c) - \hat{p}_i(c))^2 \end{aligned}$$

# 其实可以这样

- 很少的数学知识
- 很少的机器学习
- 基本的Python

# 选择正确的参考书



# Framework



TensorFlow



PyTorch




MindSpore

Jittor



# Lane Segmentation





# 项目设计原则


- 紧扣前沿热门领域：无人驾驶
- 既重视特定项目实战
- 又重视算法的普适性
- 关注领域最新发展

Make  
your  
hand  
dirty



Howard Chow

✌️ 第一个完整手撸的工程。深度神经网络真的太神奇了，它咋知道那就是车道线呀😂



1小时前



♡ Alan Wang, AI-King

Rongfan Leo: 效果还挺好的

# Lane Segmentation

- 骨干网络：ResNet
- 语义分割进入深度学习时代：FCN
- 简单实用：U-Net
- Google出品：Deeplab v3+

# Lane Segmentation

- 项目框架实战
- 模型训练实战
- 模型部署实战
- 图像分割最新发展：全景分割

# Notes

# Backbone Network

- VGG
- ResNet

# VGG



# VGG

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input ( $224 \times 224$ RGB image)					
conv3-64	conv3-64 <b>LRN</b>	conv3-64 <b>conv3-64</b>	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 <b>conv3-128</b>	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 <b>conv1-256</b>	conv3-256 conv3-256 <b>conv3-256</b>	conv3-256 conv3-256 conv3-256 <b>conv3-256</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					




# VGG-BN

```
def make_layers(cfg, batch_norm=False):
    layers = []
    in_channels = 3
    for v in cfg:
        if v == 'M':
            layers += [nn.MaxPool2d(kernel_size=2, stride=2)]
        else:
            conv2d = nn.Conv2d(in_channels, v, kernel_size=3, padding=1)
            if batch_norm:
                layers += [conv2d, nn.BatchNorm2d(v), nn.ReLU(inplace=True)]
            else:
                layers += [conv2d, nn.ReLU(inplace=True)]
            in_channels = v
    return nn.Sequential(*layers)
```



# 参数初始化

```
def _initialize_weights(self):
    for m in self.modules():
        if isinstance(m, nn.Conv2d):
            nn.init.kaiming_normal_(m.weight, mode='fan_out', nonlinearity='relu')
            if m.bias is not None:
                nn.init.constant_(m.bias, 0)
        elif isinstance(m, nn.BatchNorm2d):
            nn.init.constant_(m.weight, 1)
            nn.init.constant_(m.bias, 0)
        elif isinstance(m, nn.Linear):
            nn.init.normal_(m.weight, 0, 0.01)
            nn.init.constant_(m.bias, 0)
```

# ResNet



# Deep Network Degradation



# Notes


# Residual Learning

$$\mathcal{H}(\mathbf{x})$$

$$\mathcal{F}(\mathbf{x}) := \mathcal{H}(\mathbf{x}) - \mathbf{x}.$$

$$\mathcal{H}(\mathbf{x}) = \mathcal{F}(\mathbf{x}) + \mathbf{x}.$$

# Residual Learning



# Notes



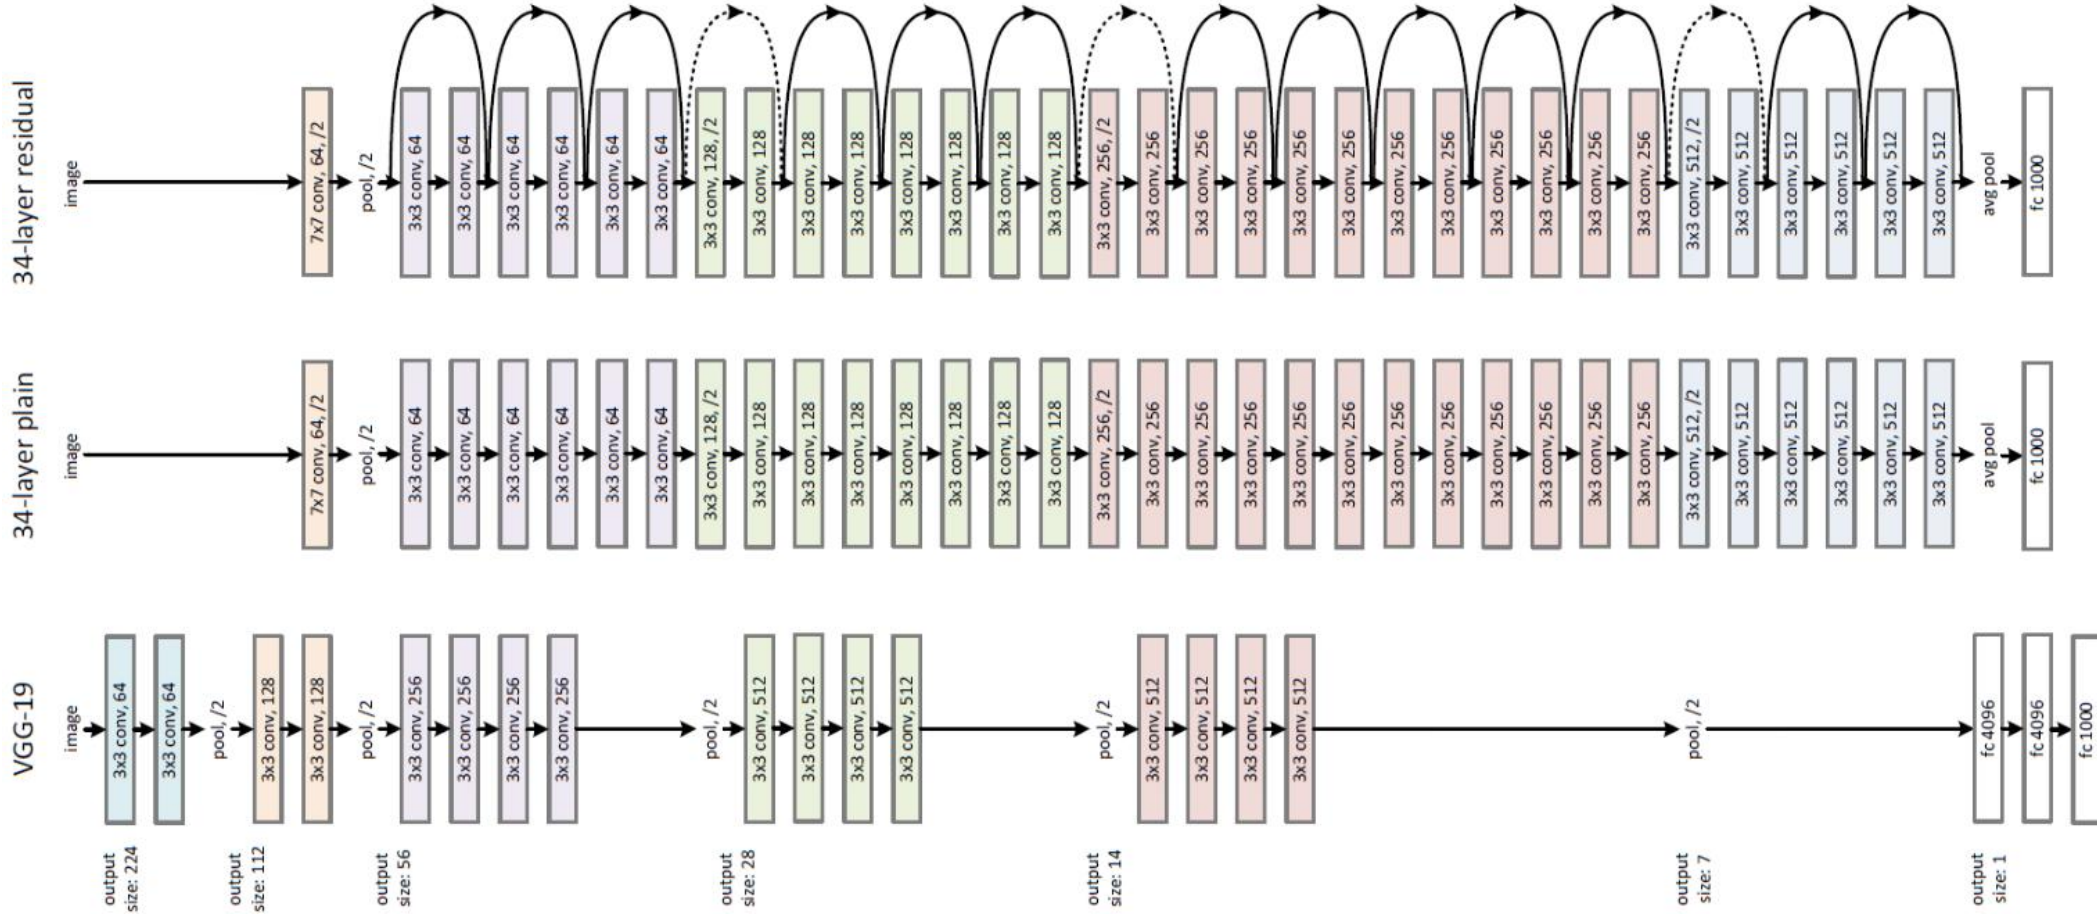
# Identity Mapping Shortcut

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{W_i\}) + \mathbf{x}.$$

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{W_i\}) + W_s \mathbf{x}.$$

- 1x1 Convolution


# ResNet



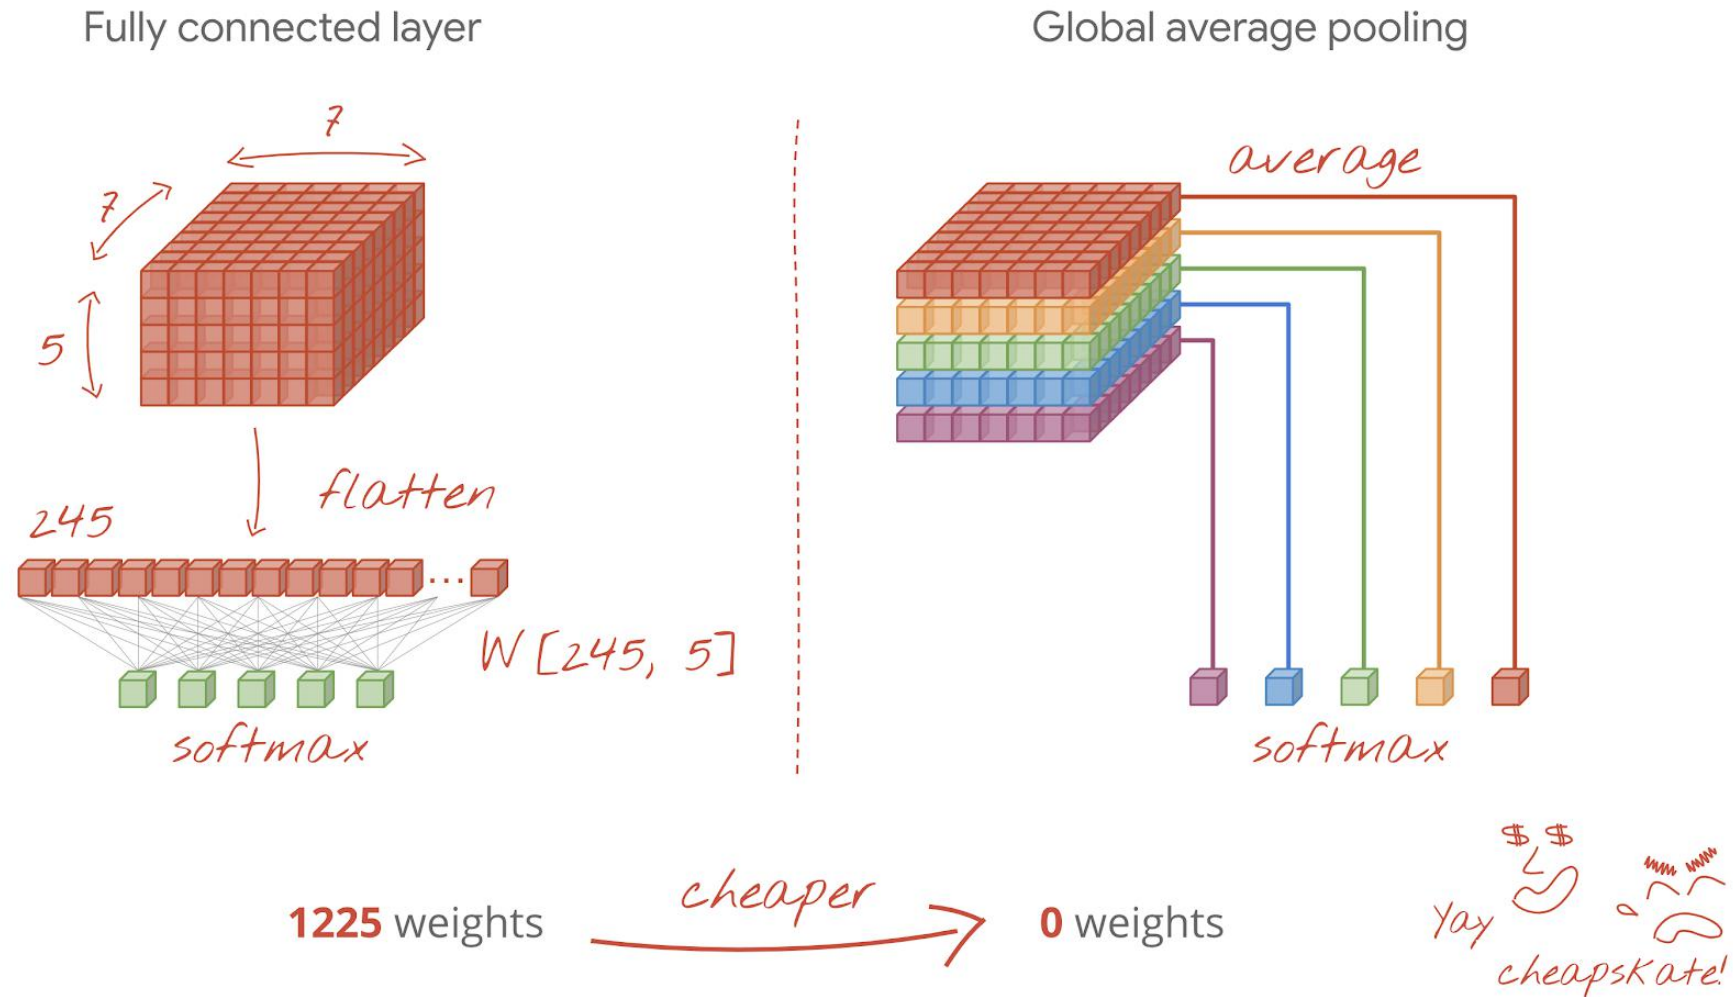
# ResNet

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	7×7, 64, stride 2				
conv2_x	56×56	3×3 max pool, stride 2				
		$\begin{bmatrix} 3\times 3, 64 \\ 3\times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 64 \\ 3\times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3\times 3, 128 \\ 3\times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 128 \\ 3\times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3\times 3, 256 \\ 3\times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 256 \\ 3\times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3\times 3, 512 \\ 3\times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 512 \\ 3\times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

# Global Average Pooling




# Global Average Pooling





# Global Average Pooling

- `torch.nn.AdaptiveAvgPool2d(output_size)`
- `tf.keras.layers.GlobalAvgPool2D`



# BottleNeck




# ResNet新发展

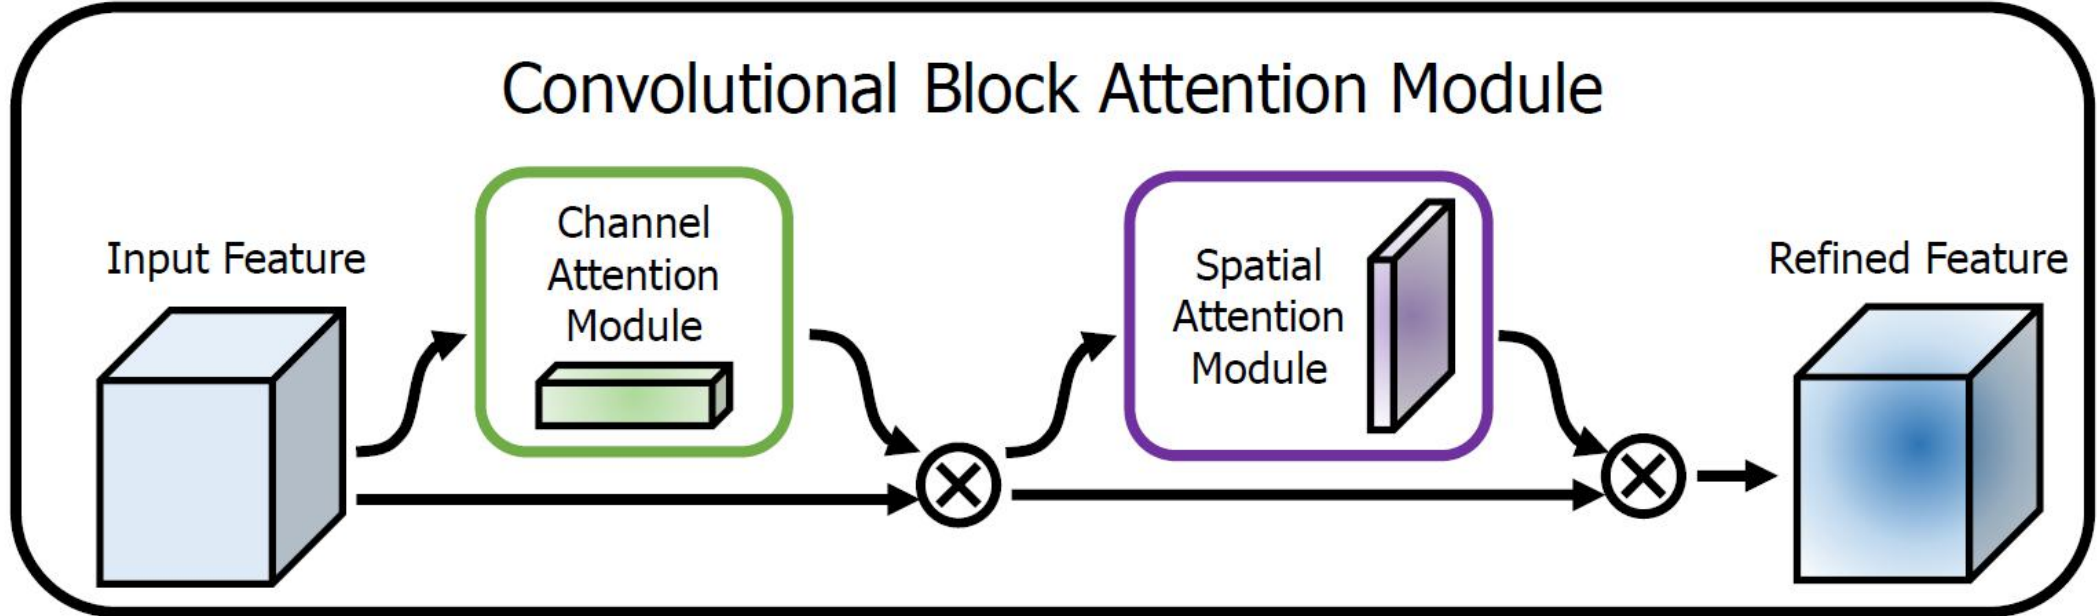
- ResNeXt
- ResNeXt-Attention
- ResNeXt WSL
- ResNeSt



# ResNeXt



# ResNeXt-Attention



# SENet




Figure 1: A Squeeze-and-Excitation block.

# ResNeXt WSL


[Get Started](#)[Ecosystem](#)[Mobile](#)[Blog](#)[Tutorials](#)[Docs](#)[Resources](#)[GitHub](#)[View on Github](#)[Open on Google Colab](#)

Fig. 5: Classification accuracy on val-IN-1k using ResNeXt-101  $32 \times \{4, 8, 16, 32, 48\}d$  with and without pretraining on the IG-940M-1.5k dataset.

```
import torch
model = torch.hub.load('facebookresearch/WSL-Images', 'resnext101_32x8d_wsl')
# or
# model = torch.hub.load('facebookresearch/WSL-Images', 'resnext101_32x16d_wsl')
# or
# model = torch.hub.load('facebookresearch/WSL-Images', 'resnext101_32x32d_wsl')
# or
# model = torch.hub.load('facebookresearch/WSL-Images', 'resnext101_32x48d_wsl')
model.eval()
```

All pre-trained models expect input images normalized in the same way, i.e. mini-batches of 3-channel RGB images of shape  $(3 \times H \times W)$ , where  $H$  and  $W$  are expected to be at least 224. The images have to be loaded in to a range of  $[0, 1]$  and then normalized using `mean = [0.485, 0.456, 0.406]` and `std = [0.229, 0.224, 0.225]`.

# ResNeSt

## ResNeSt: Split-Attention Networks

Hang Zhang, Chongruo Wu\*, Zhongyue Zhang, Yi Zhu, Haibin Lin, Zhi Zhang,  
Yue Sun, Tong He, Jonas Mueller, R. Manmatha, Mu Li, and Alexander Smola

Amazon, University of California, Davis\*

{hzaws, chongrwu, zhongyue, yzaws, haibilin, zhiz, ysunmzn,  
htong, jonasmue, manmatha, mli, smola}@amazon.com

**Abstract.** While image classification models have recently continued to advance, most downstream applications such as object detection and semantic segmentation still employ ResNet variants as the backbone net-

[cs.CV] 16 Apr 2020

# Notes

# 课程总结

- 计算机视觉学习提速绝招
- 项目概述
- ResNet
- ResNeXt

# 重难点

- ResNet



# 参考资料

- 动手学深度学习
- Very deep convolutional networks for large-scale image recognition
- Deep Residual Learning for Image Recognition
- Aggregated Residual Transformations for Deep Neural Networks

ResNeSt: Split-Attention Networks

# 课程作业

- 使用Pytorch逐行实现ResNet-152
- 加载Pytorch官方预训练参数(ImageNet-1000)
- 实现图像识别功能

# Week 2: FCN

